

More programs, or more rolls?

386 MATH-500 tasks · 9 slots per arm · 3 repeats per slot

Generated panel



8 generated programs + baseline

vs

Same-code control



9 byte-identical baseline slots

a Coverage

Repeat-mean oracle headroom

Generated



Same code



0 2.6 pp

Identical code still produces positive oracle headroom.

b Repeatability

Residual score correlation

Generated

0.789

Same code

0.006



All 3 repeats vs baseline:

100 loss tasks · 1 win task*

c Usable selection

Frozen selector gain

0.00 pp

vs development-fixed baseline

27-execution oracle coverage

98.70%

Equal in both arms

Clone-adjusted held-out gain: -0.26 pp · 95% CI $[-1.45, +0.95]$

Stable complementarity remains unresolved.

Error bars: 95% bootstrap CIs. *The single win is sensitive to answer extraction.

Matched harness executions; model-call and token costs vary. Program icons are schematic.