

Generated



8 programs + baseline

386 common MATH-500 tasks · 3 repeats/slot

Same code



9 identical baseline slots

Discover → freeze → validate




3 folds

2 repeats: freeze choices
per-task member + best fixed

Held-out repeat: gain G

(a) Coverage

Oracle · execution-matched



Repeat-mean oracle headroom · pp



(b) Repeatability

Residual r · 95% CI




Persistent scores vs baseline

Any member · 3/3 repeats · tasks



(c) Useful selection

Clone-adjusted held-out gain · pp



Complementarity unresolved

Frozen feature selector

