

386 complete MATH-500 tasks · 3 repeats/slot

● **8 generated**
+ baseline



■ **9 identical**
baseline slots

3-fold: discover → freeze → test




2 repeats: freeze choices
per-task + best fixed member

Held-out repeat: gain G

(a) Coverage

Oracle · execution-matched



Repeat-mean headroom (pp)



(b) Repeatability

Residual r · 95% CI




Persistent scores vs baseline

Tasks · any member · 3/3 repeats



(c) Useful selection

Clone-adjusted held-out gain (pp)



Specialization unresolved

📁 **Frozen feature selector** gain
→ **Baseline** (386/386) **0.00 pp**