

V8 · A — calibrated




Figure 1: Coverage, repeatable scores, and useful selection on 386 complete MATH tasks. Both arms freeze per-task choices and a best fixed member on two repeats, then measure gain G on the held-out repeat. D subtracts the same-code gain. Oracle budgets count harness executions. Intervals are 95% bootstrap CIs; the sole persistent win is extraction-sensitive. Persistent task counts require at least one generated member to differ from baseline in all three repeats.

The evidence yields three contributions. **First**, identical programs produce positive repeat-averaged oracle headroom, so that signal alone cannot identify program specialization. **Second**, generated programs have substantially more repeatable scored-outcome patterns than clones, dominated by persistent losses: losses span 100 tasks (570 member–task pairs), whereas persistent wins occur on one task. Stable differences therefore need not provide useful task advantages; the clone-adjusted held-out gain is -0.26 pp. **Third**, the frozen selector gains 0.00 pp, and the 27-execution oracle replay ties at 98.70%. On this panel, more programs produce neither a gain for the tested selector nor greater full-budget oracle coverage than more executions of the baseline.

2 RELATED WORK

Automated agent and harness design. ADAS treats agent programs as a search space (Hu et al., 2024). Self-Harness mines failures from traces, proposes targeted edits, and retains edits after regression testing (Zhang et al., 2026b); TTHE adapts the executable harness during evaluation using unlabeled traces (Nie et al., 2026). Other systems improve task adaptation and validation allocation (Zhang et al., 2026a; Miyai et al., 2026), or make harness behavior observable and editable through verification, diagnosis, localization, and policy evolution (Xu et al., 2026; Chen et al., 2026a; Lin et al., 2026; Luo et al., 2026; Wang et al., 2026; Chen et al., 2026b). El et al. (2025) find low diversity in task-level score patterns among the automatically designed agents they study. We separate two explanations for population-level results: distinct programs can execute the same behavior, and identical programs can produce different outcomes. Trace inspection tests whether proposed mechanisms execute; repeated same-code controls isolate the coverage available without changing the program. This connects program equivalence and mutation testing (Budd & Angluin, 1982; Papadakis et al., 2019) to performance attribution in generated harness populations.

V8 · B — fingerprint




Figure 1: Coverage, repeatable scores, and useful selection on 386 complete MATH tasks. Both arms freeze per-task choices and a best fixed member on two repeats, then measure gain G on the held-out repeat. D subtracts the same-code gain. Oracle budgets count harness executions. Intervals are 95% bootstrap CIs; the sole persistent win is extraction-sensitive. Example #138 is deliberately selected to illustrate the recorded outcomes; it is not representative.

The evidence yields three contributions. **First**, identical programs produce positive repeat-averaged oracle headroom, so that signal alone cannot identify program specialization. **Second**, generated programs have substantially more repeatable scored-outcome patterns than clones, dominated by persistent losses: losses span 100 tasks (570 member–task pairs), whereas persistent wins occur on one task. Stable differences therefore need not provide useful task advantages; the clone-adjusted held-out gain is -0.26 pp. **Third**, the frozen selector gains 0.00 pp, and the 27-execution oracle replay ties at 98.70%. On this panel, more programs produce neither a gain for the tested selector nor greater full-budget oracle coverage than more executions of the baseline.

2 RELATED WORK

Automated agent and harness design. ADAS treats agent programs as a search space (Hu et al., 2024). Self-Harness mines failures from traces, proposes targeted edits, and retains edits after regression testing (Zhang et al., 2026b); TTHE adapts the executable harness during evaluation using unlabeled traces (Nie et al., 2026). Other systems improve task adaptation and validation allocation (Zhang et al., 2026a; Miyai et al., 2026), or make harness behavior observable and editable through verification, diagnosis, localization, and policy evolution (Xu et al., 2026; Chen et al., 2026a; Lin et al., 2026; Luo et al., 2026; Wang et al., 2026; Chen et al., 2026b). El et al. (2025) find low diversity in task-level score patterns among the automatically designed agents they study. We separate two explanations for population-level results: distinct programs can execute the same behavior, and identical programs can produce different outcomes. Trace inspection tests whether proposed mechanisms execute; repeated same-code controls isolate the coverage available without changing the program. This connects program equivalence and mutation testing (Budd & Angluin, 1982; Papadakis et al., 2019) to performance attribution in generated harness populations.

V8 · C — crossover




Figure 1: Coverage, repeatable scores, and useful selection on 386 complete MATH tasks. Both arms freeze per-task choices and a best fixed member on two repeats, then measure gain G on the held-out repeat. D subtracts the same-code gain. Oracle budgets count harness executions. Intervals are 95% bootstrap CIs; the sole persistent win is extraction-sensitive. The top-right sketches are population success-probability examples, not observed task curves. Stable complementarity requires the better member to change across tasks.

The evidence yields three contributions. **First**, identical programs produce positive repeat-averaged oracle headroom, so that signal alone cannot identify program specialization. **Second**, generated programs have substantially more repeatable scored-outcome patterns than clones, dominated by persistent losses: losses span 100 tasks (570 member–task pairs), whereas persistent wins occur on one task. Stable differences therefore need not provide useful task advantages; the clone-adjusted held-out gain is -0.26 pp. **Third**, the frozen selector gains 0.00 pp, and the 27-execution oracle replay ties at 98.70%. On this panel, more programs produce neither a gain for the tested selector nor greater full-budget oracle coverage than more executions of the baseline.

2 RELATED WORK

Automated agent and harness design. ADAS treats agent programs as a search space (Hu et al., 2024). Self-Harness mines failures from traces, proposes targeted edits, and retains edits after regression testing (Zhang et al., 2026b); TTHE adapts the executable harness during evaluation using unlabeled traces (Nie et al., 2026). Other systems improve task adaptation and validation allocation (Zhang et al., 2026a; Miyai et al., 2026), or make harness behavior observable and editable through verification, diagnosis, localization, and policy evolution (Xu et al., 2026; Chen et al., 2026a; Lin et al., 2026; Luo et al., 2026; Wang et al., 2026; Chen et al., 2026b). El et al. (2025) find low diversity in task-level score patterns among the automatically designed agents they study. We separate two explanations for population-level results: distinct programs can execute the same behavior, and identical programs can produce different outcomes. Trace inspection tests whether proposed mechanisms execute; repeated same-code controls isolate the coverage available without changing the program. This connects program equivalence and mutation testing (Budd & Angluin, 1982; Papadakis et al., 2019) to performance attribution in generated harness populations.